

NYUv2 surface normals — d_l (%) ↑

Method	DINOv3-S	DINOv3-B	DINOv3-L	MAE-L	SigLIP2-L	V-JEPA2.1-L	PESpatial-L	Params (M)	Rank
LastLayer	22.31	24.52	23.50	32.81	13.22	36.42	16.88	1.05	4.57
AvgLayer	20.74	27.89	30.99	35.32	15.17	32.63	17.42	1.05	4.29
WeightedLayers	25.44	30.80	34.65	36.58	16.74	35.52	17.46	1.05	3.14
ConcatLayers	40.20	45.19	48.48	49.87	36.12	49.10	33.41	26.22	1.00
AMDF (ours)	<u>37.31</u>	<u>41.98</u>	<u>45.55</u>	<u>46.90</u>	<u>32.61</u>	<u>46.34</u>	<u>29.58</u>	1.17	2.00

NYUv2 depth — d_l (%) ↑ (log-depth head)

Method	DINOv3-S	DINOv3-B	DINOv3-L	MAE-L	SigLIP2-L	V-JEPA2.1-L	PESpatial-L	Params (M)	Rank
LastLayer	58.80	61.73	63.17	58.10	55.16	<u>67.63</u>	54.11	1.05	3.71
AvgLayer	51.90	60.07	60.75	59.30	57.51	64.46	54.93	1.05	4.57
WeightedLayers	56.13	61.34	62.15	59.77	57.55	66.07	55.13	1.05	3.57
ConcatLayers	61.13	64.86	67.35	<u>64.11</u>	<u>62.28</u>	67.44	<u>57.62</u>	26.22	1.71
AMDF (ours)	<u>58.81</u>	<u>64.06</u>	<u>66.47</u>	64.22	63.01	69.80	58.56	1.16	1.43

NYUv2 semantic segmentation — mIoU (%) ↑

Method	DINOv3-S	DINOv3-B	DINOv3-L	MAE-L	SigLIP2-L	V-JEPA2.1-L	PESpatial-L	Params (M)	Rank
LastLayer	44.20	53.41	<u>57.08</u>	38.94	40.39	44.05	32.79	1.09	3.43
AvgLayer	35.49	47.48	50.80	40.54	42.13	41.94	33.36	1.09	4.29
WeightedLayers	40.26	50.02	52.78	41.00	42.11	42.94	33.25	1.09	3.86
ConcatLayers	44.61	52.63	56.50	43.89	47.07	46.26	36.48	26.26	1.57
AMDF (ours)	<u>44.34</u>	<u>53.13</u>	57.17	<u>42.36</u>	<u>43.83</u>	<u>44.77</u>	<u>33.64</u>	1.20	1.86

Taskonomy — aggregate $L_l \times 100$ ↓

Method	DINOv3-S	DINOv3-B	DINOv3-L	MAE-L	SigLIP2-L	V-JEPA2.1-L	PESpatial-L	Params (M)	Rank
LastLayer	41.64	39.21	39.26	39.50	46.81	36.94	42.21	1.06	4.43
AvgLayer	43.66	39.45	38.51	38.60	45.48	37.84	41.83	1.06	4.43
WeightedLayers	41.47	38.51	37.51	38.19	44.30	37.32	41.70	1.06	3.14
ConcatLayers	33.90	32.19	30.95	32.67	35.94	31.40	34.79	26.23	1.00
AMDF (ours)	<u>37.49</u>	<u>33.86</u>	<u>32.21</u>	<u>33.77</u>	<u>37.05</u>	<u>34.78</u>	<u>36.83</u>	1.17	2.00

NAVI camera pose — rotation (°) ↓ / translation ↓

Method	DINOv3-S		DINOv3-B		DINOv3-L		MAE-L		SigLIP2-L		V-JEPA2.1-L		PESpatial-L	
	Rot.	Trans.	Rot.	Trans.	Rot.	Trans.	Rot.	Trans.	Rot.	Trans.	Rot.	Trans.	Rot.	Trans.
LastLayer	45.86	0.17	44.24	0.17	40.63	0.16	38.61	0.16	62.62	0.22	<u>37.34</u>	0.15	65.16	0.23
AvgLayer	44.76	0.17	36.04	0.15	36.44	0.14	37.28	0.15	59.01	0.21	47.82	0.19	66.44	0.24
WeightedLayers	43.71	0.17	35.25	0.15	35.09	0.14	35.82	0.15	58.60	0.21	43.59	<u>0.17</u>	66.21	0.23
ConcatLayers	34.01	0.14	26.55	0.12	24.31	0.11	31.10	0.13	32.29	0.14	37.46	0.15	52.28	0.19
AMDF (ours)	<u>40.90</u>	<u>0.16</u>	<u>31.61</u>	<u>0.14</u>	<u>31.80</u>	<u>0.13</u>	<u>31.70</u>	<u>0.14</u>	<u>47.37</u>	<u>0.18</u>	36.72	0.15	<u>63.77</u>	<u>0.22</u>

Bold red = best, underlined = second best per column (among the 5 rows). Rank = mean over backbone columns of the per-column rank (ties averaged; NAVI over 14 rot/trans sub-columns). Params (M) = trainable parameters (probe + task head), DINOv3-L. Full-width setup: no channel reduction, 64-d depth keys/queries, per-layer expert queries (AMDF), 16 experts; frozen backbones; same recipe for every row of a table; mean of up to 3 seeds (8/42/123; a few baseline cells have 1–2), best-epoch validation metric. Depth uses a log-depth head.